



Argonne
NATIONAL
LABORATORY

... for a brighter future



U.S. Department
of Energy

UChicago ►
Argonne_{LLC}



**Office of
Science**
U.S. DEPARTMENT OF ENERGY

A U.S. Department of Energy laboratory
managed by UChicago Argonne, LLC

MPITypes: **Processing MPI Datatypes** **Outside MPI**

**Rob Ross¹, Rob Latham¹, William Gropp²,
Ewing Lusk¹, Rajeev Thakur¹**

**¹ Mathematics and Computer Science Division
Argonne National Laboratory
{rross, robl, lusk, thakur}@mcs.anl.gov**

**² Computer Science Department
University of Illinois at Urbana-Champaign
wgropp@illinois.edu**

Motivation – Libraries and MPI

- Libraries for parallel computing play a critical role in improving the performance of codes and productivity of application writers (e.g., MPI libraries, ScaLAPACK, PETSc, HDF5)
- MPI communicators, requests, attributes, and datatypes are extremely useful constructs for building parallel libraries
- Some improvements could be made in generalized requests
 - R. Latham, W. Gropp, R. Ross, and R. Thakur, “Extending the MPI-2 Generalized Request Interface,” Proc. of EuroPVM/MPI 2007.

The biggest missing piece for parallel libraries that build on MPI is a system for efficient, custom manipulation of data described by MPI datatypes.



Custom MPI Datatype Processing

We need more than MPI_Pack and MPI_Unpack.

- ROMIO – data sieving and two-phase optimizations
 - Operates on portions of types (partial processing)
 - Combines multiple types together
 - Types describing file regions, not just memory regions
- Parallel netCDF
 - May operate on portions of types
 - Byteswaps data on some systems
 - Converts data from one representation to another

The MPITypes Library

MPITypes is a portable, open source library for processing MPI datatypes in libraries and applications.

- Based on MPICH2 datatype processing component
- Built-in functions for packing, unpacking, and flattening
- Toolkit for building custom type processing routines
- Uses only MPI-2 functionality for accessing datatype information and caching data:
 - Datatype envelope and contents functions
 - Attributes on communicators and datatypes

Outline of Talk

- Motivation
- **Datatype processing in MPICH2**
 - **Dataloop representation of MPI datatypes**
 - **Segments, leaf functions, and traversing dataloop trees**
- MPITypes
 - Summary
 - Basic functionality
 - Building functions with MPITypes
- Performance evaluation
- Related work
- Concluding remarks



Datatype Processing (from MPICH2)

- **Uses a simplified representation, called dataloops**
- Five basic dataloop node types with increasing complexity
 - **contig** - blocklength
 - **vector** - count, blocklength, stride
 - **blockindexed** - count, blocklength, offset array
 - **indexed** - count, blocklength array, offset array
 - **struct** - count, blocklength array, offset array
- Nodes are used to build trees for more complex types
- Dataloop tree is processed nonrecursively, with state held in a data structure called a **segment**
- **Leaf functions** define the operation performed on data



Simple Dataloop Example

- MPI Vector datatype:

MPI_Vector, cnt = 2, blklen = 2, str = 4

MPI_INT



- Dataloop representation (just one leaf node, 8-byte integers):

DLP_Vector, cnt = 2, blklen = 2, str = 32,
el_sz = 8, el_ext = 8, el_type = MPI_INT



Complex Dataloop Example (FLASH)



■ Dataloop representation:

DLP_Vector, cnt = 80, blklen = 1, str = 768432,
el_sz = 4096, el_ext = 366920

DLP_Vector, cnt = 8, blklen = 1, str = 49152,
el_sz = 512, el_ext = 22856


DLP_Vector, cnt = 8, blklen = 1, str = 3072,
el_sz = 64, el_ext = 1352

DLP_Vector, cnt = 8, blklen = 1, str = 192,
el_sz = 8, el_ext = 8,
el_type = MPI_DOUBLE



Traversing Dataloops

```
if (not a leaf node) {  
    while (not done with this dataloop node) {  
        update segment with new position  
        push current dataloop state onto stack  
        process next dataloop in tree  
        decrement count/blklen in segment  
    }  
    pop dataloop off the stack and resume processing  
} else /* leaf */ {  
    if (leaf type is index && have index leaf fn) call index leaf fn  
    if (leaf type is vector && have vector leaf fn) call vector leaf fn  
    ...  
    else call contig leaf fn  
    pop dataloop off stack in segment and resume processing  
}
```



Outline of Talk

- Motivation
- Datatype processing in MPICH2
 - Dataloop representation of MPI datatypes
 - Segments, leaf functions, and traversing dataloop trees
- **MPITypes**
 - **Summary**
 - **Basic functionality**
 - **Building functions with MPITypes**
- Performance evaluation
- Related work
- Concluding remarks



The MPITypes Library

- Datatype processing extracted from MPICH2
 - Symbols renamed to avoid conflicts
 - Builds using mpicc of local MPI implementation
- MPITypes functions build a dataloop representation on demand
 - `MPI_Type_get_envelope` and `MPI_Type_get_contents` used to extract datatype parameters
 - Attribute on datatype is used to associate the dataloop representation with the type, so it is only built once
- Segment structure used to maintain state during processing (as in MPICH2)
- User-defined processing functions allow custom behavior



Basic MPITypes Operators

int **MPIT_Type_memcpy**(typebuf, count, type, streambuf, direction, start, end);

- *Like MPI_Pack/MPI_Unpack, but allows partial processing*

int **MPIT_Type_blockct**(count, type, start, end, blockct);

- *Provides a count of the number of contiguous regions in a portion of a [count, datatype] tuple*

int **MPIT_Type_flatten**(typebuf, count, type, start, end, disps, blocklens, count);

- *Generates a list of offsets and lengths for a portion of a [buffer, count, datatype] tuple*

MPITypes Toolkit Functions

MPIT_Segment ***MPIT_Segment_alloc**();

int **MPIT_Segment_init**(buf, count, type, segment, flag);

int **MPIT_Segment_free**(segment);

- *Allocate, initialize, and free the data structure used to track progress while processing a [count, datatype] tuple*

int **MPIT_Segment_manipulate**(segment, start, end,
(*contigfn) (...), (*vectorfn) (...), (*blkidxfn) (...), (*indexfn) (...),
(*sizefn) (el_type), params);

- *Traverse the datatype representation, executing specified leaf functions as leaves are encountered*

Basic MPITypes operators are built with these functions.

MPITypes memcpy Implementation (1/2)

```
typedef struct MPIT_memcpy_params_s {  
    int direction; char *packbuf, *userbuf;  
} MPIT_memcpy_params;
```

**Data structure used to
hold parameters relevant
to custom processing**

```
int MPIT_Leaf_contig_memcpy(MPI_Aint *blocks_p, MPI_Type el_type,  
    MPI_Aint dtype_pos, void *unused, void *v_paramp)  
{  
    MPI_Aint size, el_size; MPIT_memcpy_params *paramp = v_paramp;  
    MPI_Type_size(el_type, &el_size); size = *blocks_p * el_size;  
  
    if (paramp->direction == MPIT_MEMCPY_TO_USERBUF)  
        memcpy(paramp->userbuf + dtype_pos, paramp->packbuf, size);  
    else  
        memcpy(paramp->packbuf, paramp->userbuf + dtype_pos, size);  
  
    paramp->packbuf += size; return 0;  
}
```



MPITypes memcpy Implementation (2/2)

```
int MPIT_Type_memcpy(void *typebuf, int count, MPI_Datatype type,
    void *streambuf, int direction, MPI_Aint start, MPI_Aint *end)
{
    MPIT_Segment *segp;
    MPIT_memcpy_params params;

    segp = MPIT_Segment_alloc();
    MPIT_Segment_init(NULL, count, type, segp, 0);

    params.userbuf  = typebuf;
    params.packbuf  = packbuf;
    params.direction = direction;

    MPIT_Segment_manipulate(segp, start, end,
        MPIT_Leaf_contig_memcpy, NULL, NULL, NULL, NULL, &params);

    MPIT_Segment_free(segp); return MPI_SUCCESS;
}
```

Optional vector,
blockindexed,
and indexed leaf
functions would
go here.



Outline of Talk

- Motivation
- Datatype processing in MPICH2
 - Dataloop representation of MPI datatypes
 - Segments, leaf functions, and traversing dataloop trees
- MPITypes
 - Summary
 - Basic functionality
 - Building functions with MPITypes
- **Performance evaluation**
- Related work
- Concluding remarks



Performance Evaluation

■ Goals

- Show that MPITypes performs on par with current MPI implementations for a variety of datatypes
- Establish feasibility of using MPITypes to implement relevant datatype operations
- Quantify benefit of using MPITypes in the context of a parallel library

■ Evaluation

- MPITypes memcpy implementation
- MPITypes implementation of *transpacking*
- Custom function for encoding data prior to I/O in Parallel netCDF



Test Environment

- 2.66 GHz Intel Xeon node
 - 8 cores
 - 16 Gbytes of main memory
 - Little endian
- Linux 2.6.27
- GNU C compiler version 4.2.4
- MPICH2 version 1.0.8p1
 - Compiled with “--enable-fast=O3”
- Open MPI version 1.3.1
 - Compiled with “CFLAGS=-O3 --disable-heterogeneous
--enable-shared=no --enable-static
--with-mpi-param-check=no”



Comparing MPI_Pack/MPI_Unpack, MPIT_Type_memcpy, and manual copy rates

MPITypes performance is essentially identical to MPI implementations.

- Copy into and then back out of a contiguous buffer, many times (provides opportunity to verify correctness)

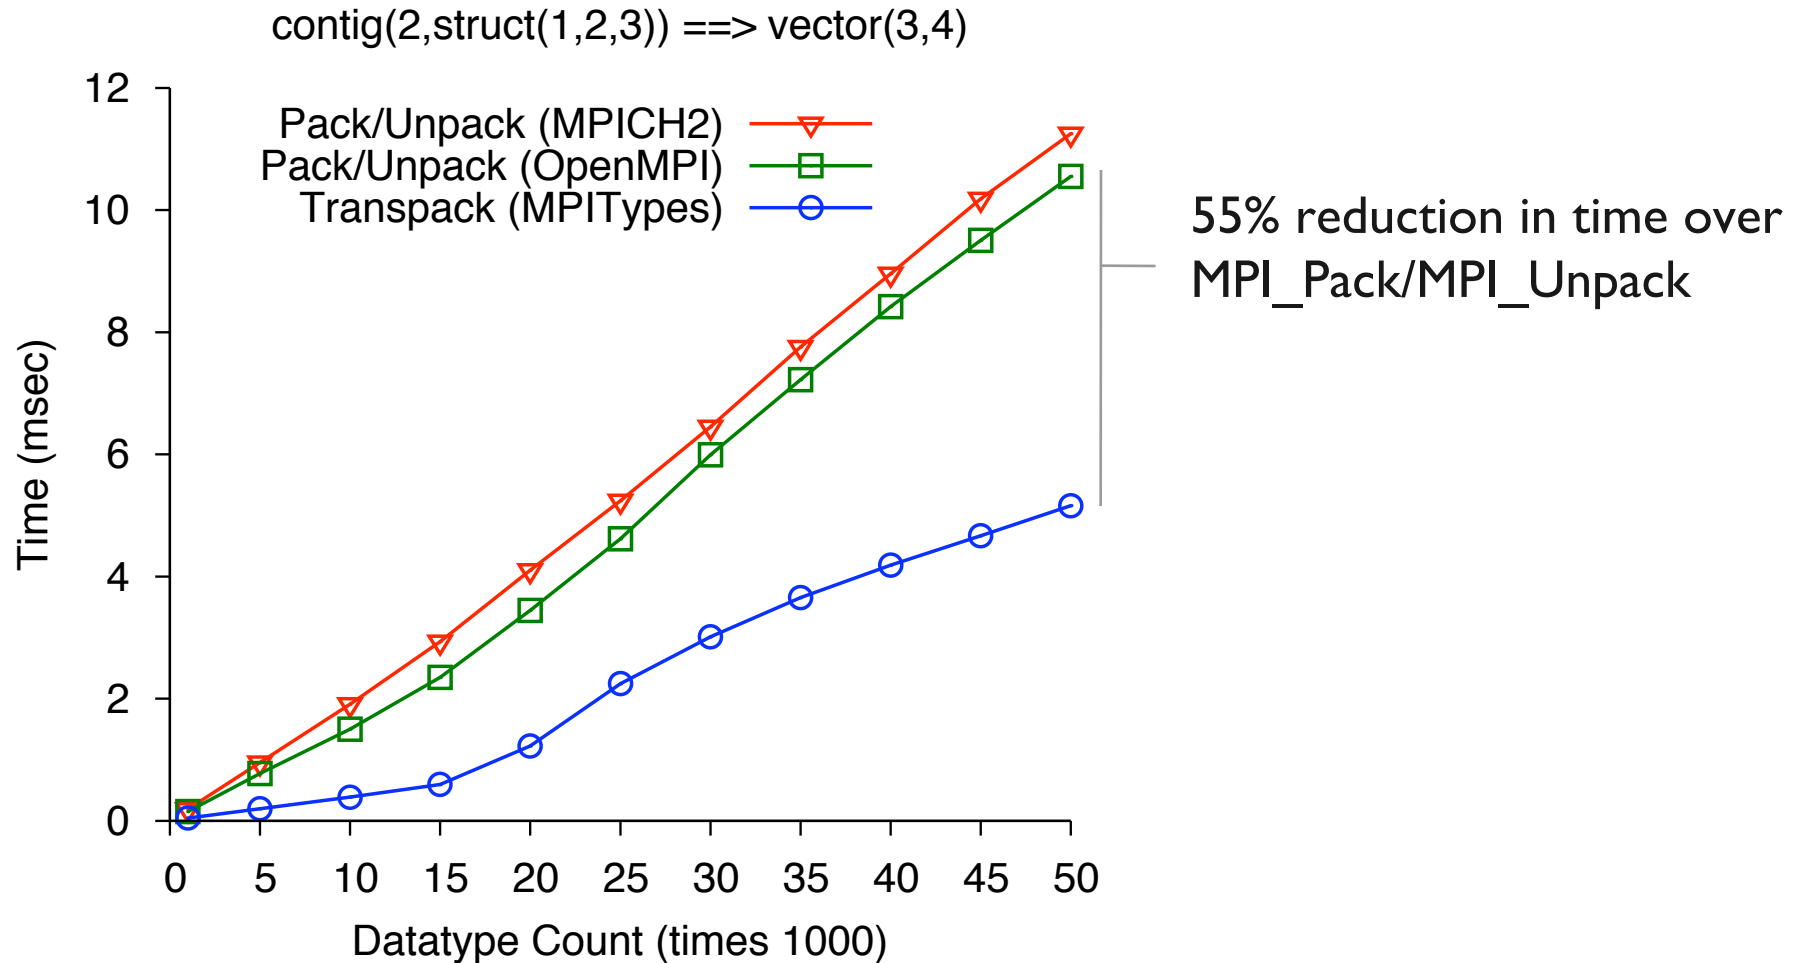
Test	MPICH2 (MB/sec)	Open MPI (MB/sec)	MPITypes (MB/sec)	Manual (MB/sec)	Size (MB)	Extent (MB)
Contig	4152.07	4157.69	4149.13	2650.81	8.00	8.00
Vector	1776.81	1680.23	1777.04	1777.60	8.00	16.00
Indexed	1120.59	967.69	1123.97	1575.41	4.00	8.00
XY Face	17564.43	18143.63	17520.11	16423.59	0.50	0.50
XZ Face	4004.26	4346.81	3975.23	3942.41	0.50	127.50
YZ Face	153.89	154.19	153.88	153.96	0.50	127.99

Transpacking


- **Transpacking is a solution to the *typed copy* problem – moving data from one datatype representation to another.**
 - Simple solution is to MPI_Pack and then MPI_Unpack, but this requires two copies and a large intermediate buffer
 - Partial processing reduces memory requirement
 - Better solution is to directly copy from one representation to another
- Quite elegant solutions to this have been proposed previously (see Mir and Träff)
- We implemented a less elegant solution using MPITypes (~200 lines)
 - Best for like-sized types with a relatively small number of contiguous regions in a single instance
 - Generates a “template” for how to copy between a single instance of each, iterates on this.



Comparing MPITypes Implementation of Transpack to MPI_Pack/MPI_Unpack (1/2)




Comparing MPITypes Implementation of Transpack to MPI_Pack/MPI_Unpack (2/2)



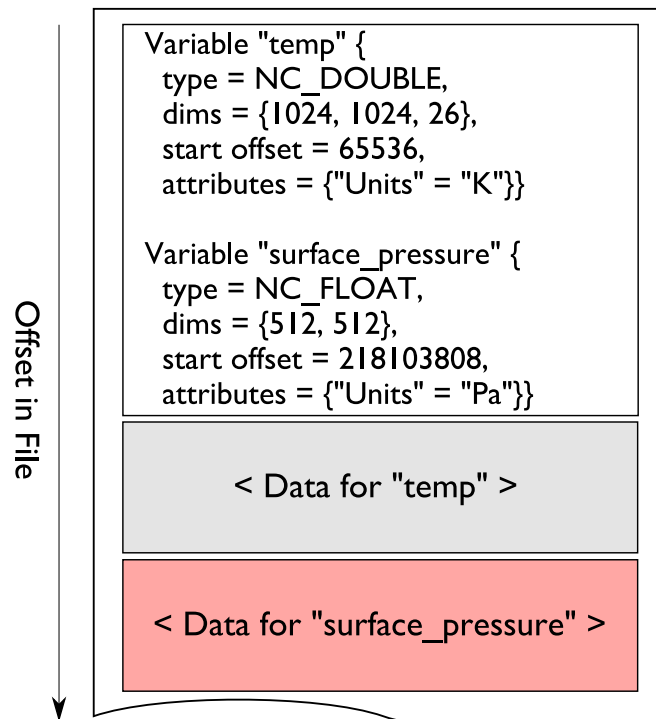
The Parallel netCDF I/O Library

Parallel netCDF (PnetCDF) provides a convenient, efficient way of storing scientific data in a portable file format.

Application Data Structures



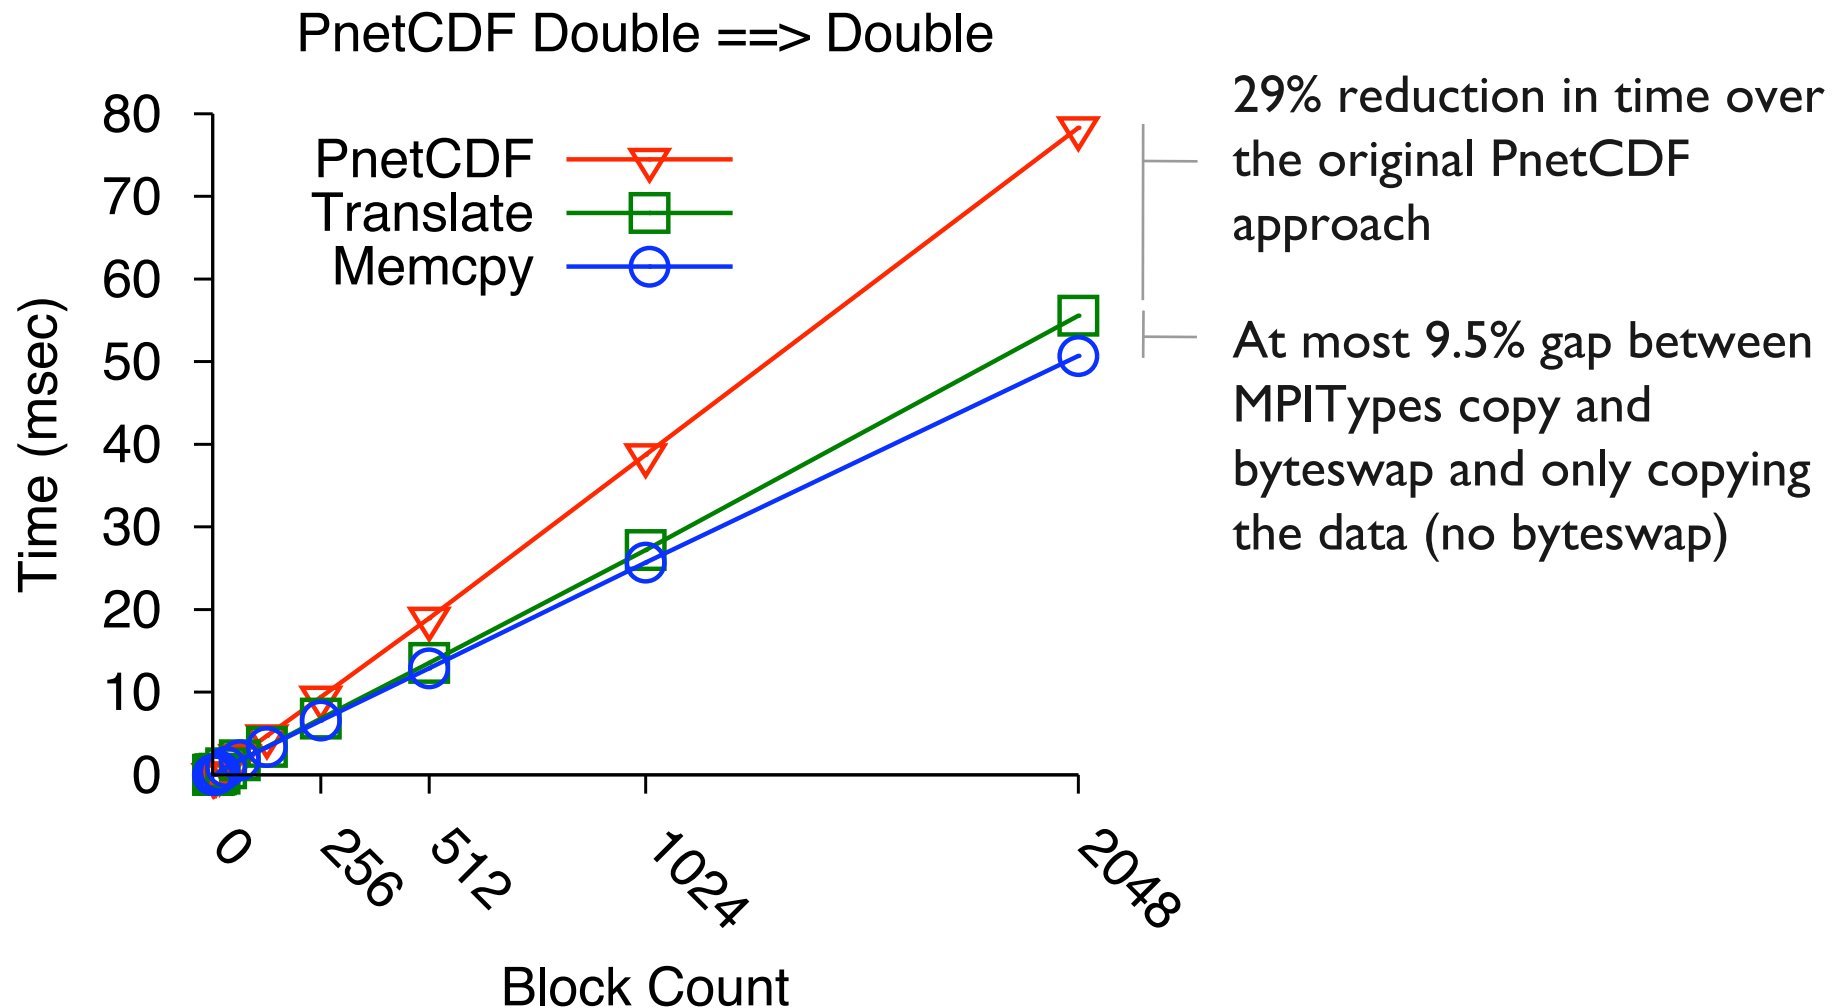
netCDF File "checkpoint07.nc"



netCDF header describes the contents of the file: typed, multi-dimensional variables and attributes on variables or the dataset itself.

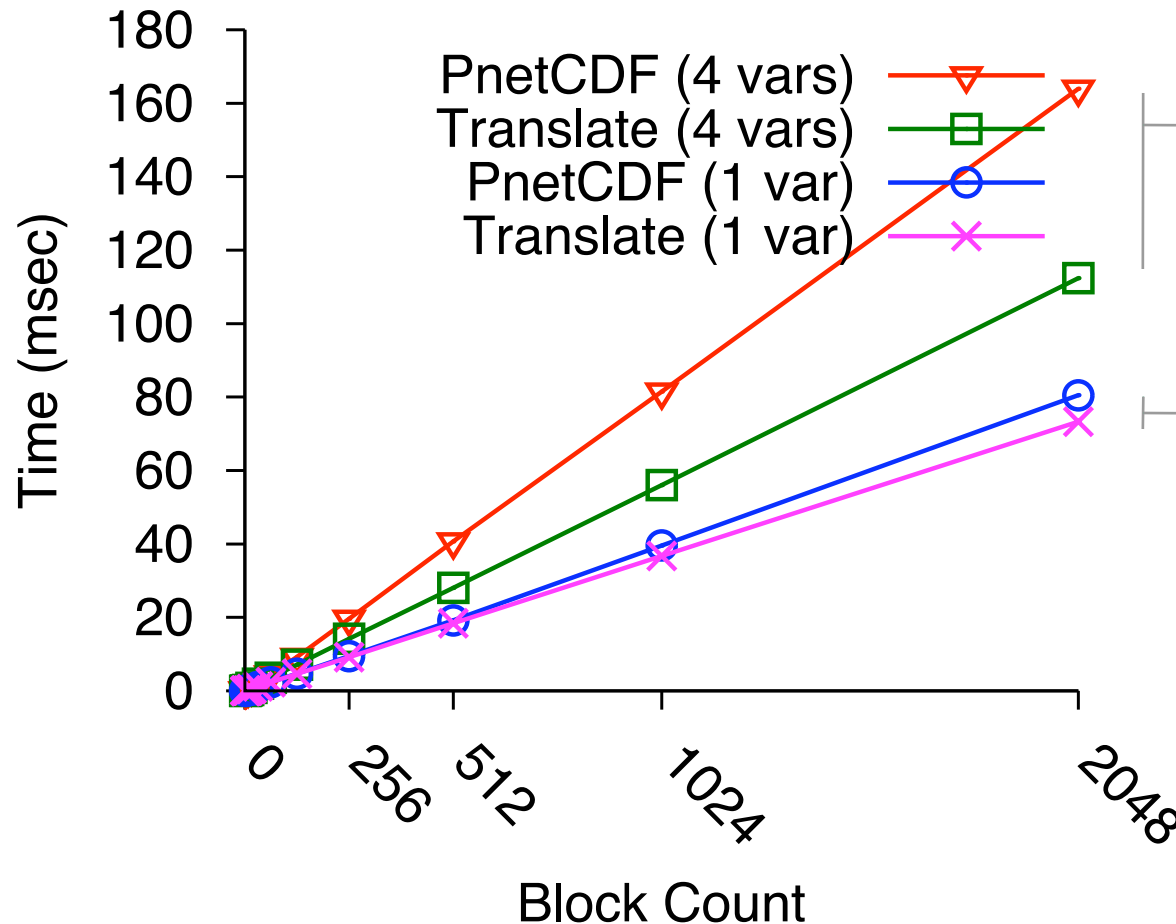
Data for variables is stored in contiguous blocks, encoded in a portable binary format according to the variable's type.

PnetCDF: Byteswapping in FLASH Case



PnetCDF: Data Conversion in FLASH Case

PnetCDF Double ==> Float



31% reduction in time over the original PnetCDF approach when four types are adjacent to one another

9% reduction in time between MPITypes and PnetCDF approach, when no elements are adjacent to one another

Related Work in Datatype Processing

- J. Träff, R. Hempel, H. Ritzdoff, and F. Zimmermann, “Flattening on the fly: Efficient handling of MPI derived datatypes,” In Proceedings of EuroPVM/MPI 1999.
- R. Ross, N. Miller, and W. Gropp, “Implementing fast and reusable datatype processing,” In Proceedings of EuroPVM/MPI 2003.
- J. Worringen, J. Träff, and H. Ritzdorf, “Improving generic noncontiguous file access for MPI-IO,” In Proceedings of EuroPVM/MPI 2003.
- F. Mir and J. Träff, “Constructing MPI input-output datatypes for efficient transpacking,” In Proceedings of EuroPVM/MPI 2008.
- F. Mir and J. Träff, “Exploiting efficient transpacking for one-sided communication and MPI-IO,” In Proceedings of EuroPVM/MPI 2009.



Concluding Remarks

- MPITypes provides a high performance, customizable implementation of datatype processing
 - Hides most of the complexity of efficiently manipulating MPI datatypes
 - Retains performance characteristics of MPI implementations
- Easy to use and incorporate into new and existing parallel libraries and applications
 - Uses MPICH2 source code license (BSD-like)
- Source code now available
 - See **<http://www.mcs.anl.gov/mpitypes>**
- Perhaps worth considering incorporating similar functionality into future MPI standards?

